

Title TBD

Subtitle TBD

UAB

**Universitat Autònoma
de Barcelona**

Niclas Rieger

Supervisor: Dr. Antonio Turiel

Dr. Álvaro Corral

Dra. Estrella Olmedo

Department of Physics

Universitat Autònoma de Barcelona

This dissertation is submitted for the degree of

Doctor of Philosophy

December 2024

Table of Contents

List of Figures	v
List of Tables	vii
Nomenclature	ix
1 Theoretical Foundations	1
1.1 Principal Component Analysis	1
1.1.1 Derivation from Maximum Variance Formulation	2
1.1.2 Sparse PCA via Rotation	7
1.1.3 Hilbert PCA	11
1.2 Maximum Covariance Analysis	14
1.3 Neural Networks	18
1.3.1 Introduction to Neural Networks	18
1.3.2 Recurrent Neural Networks	18
1.3.3 Explainable AI Techniques	19
1.4 Gaussian Processes	19
1.4.1 Introduction to Gaussian Processes	20
1.4.2 Gaussian Processes for Regression	20
1.4.3 Spatial Modeling with Gaussian Processes	20
1.4.4 Model Implementation and Inference	20
1.5 Comparative Perspective	21
2 Bridging Theory and Scientific Practice in a Data-Driven World	23
2.1 Scalable Open-Source Tools for Climate Data Analysis	23
2.2 Technical Barriers in Large-Scale and Multi-Dimensional Data Analysis	25

2.3	<i>xeofs</i> : Extracting Meaningful Patterns from Climate Data	26
	References	31
	Appendix A Proofs	39

List of Figures

1.1	Comparison of the first three modes obtained from PCA vs. Varimax-rotated PCA for a synthetic data set. (Left) Original data set with three different, uncorrelated autoregressive processes. (Middle) PCA solution. (Right) Varimax-rotated PCA solution.	12
2.1	Evolution of the OSS ecosystem.	24
2.2	Performance of Python library xeofs.	29

List of Tables

1.1	Computational complexity of PCA algorithms	7
2.1	Even better looking table using booktabs	28

Nomenclature

Greek Symbols

π $\simeq 3.14\dots$

Acronyms / Abbreviations

AMO Atlantic Multidecadal Oscillation

BSISO Boreal Summer Intraseasonal Oscillation

ENSO El Niño–Southern Oscillation

MJO Madden-Julian Oscillation

NAO North Atlantic Oscillation

PC Principal Component

PDO Pacific Decadal Oscillation

VC Vector of Coefficients

Theoretical Foundations

nicrie: Motivation for Dimensionality Reduction: Explain why dimensionality reduction is necessary in climate science. e.g. ERA5 contains more than 9 PB of data. Discuss the need to simplify complex datasets by identifying the most significant patterns and features, which can help in understanding underlying processes and improving predictability.

1.1 Principal Component Analysis

Principal Component Analysis (PCA) is a statistical technique used to reduce the dimensionality of a dataset while preserving most of the variance. This is achieved by identifying the principal axes, which are the directions in which the data varies the most. These directions correspond to the eigenvectors of the covariance matrix of the data, with the associated eigenvalues representing the magnitude of variance along these directions. PCA is particularly useful for dimensionality reduction, data compression, feature extraction and visualization of high-dimensional datasets.

The origins of PCA can be traced back to Pearson [87] in 1901 and Hotelling [52] in 1936. It gained significant traction in meteorology through the work of Lorenz [74] in 1956, who introduced the technique under the term *empirical orthogonal function* (EOF) analysis¹. The method has since been extensively applied in meteorology and climate science to identify dominant modes of variability in the atmosphere and oceans, as discussed in numerous textbooks [e.g. 28, 43, 60, 89, 101, 114].

In this section, we derive PCA from the maximum variance formulation², as outlined by Bishop [6, pp. 561-563]. This approach involves projecting the data onto a lower-dimensional linear space, known as the principal subspace, to maximize the variance of the projected data. We will then explore the limitations of PCA, particularly in the context of climate science, and discuss common strategies to address these limitations.

¹Although EOF analysis remains the preferred term in meteorology and climate science, we will use the more commonly recognized term, PCA.

²Hotelling [52] originally proposed this formulation. An alternative derivation, known as the minimum error formulation by Pearson [87], seeks a linear projection that minimizes the average projection error, defined as the mean squared distance between the original data points and their projections. This alternative approach shows that PCA is an optimal solution for approximating data in the squared error sense.

nicrie: make link to Karhunen-Loève expansion

1.1.1 Derivation from Maximum Variance Formulation

Consider a data set represented as a complex-valued, two-dimensional $N \times D$ matrix $\mathbf{Z}$, where N denotes the number of observations (or samples) and D the number of variables (or features). We will refer to the n^{th} observation as the D -dimensional vector $\mathbf{z}_n$. For simplicity, we assume the data is centered, meaning the mean of each variable is zero, i.e. $\bar{\mathbf{z}} = \frac{1}{N} \sum_{n=1}^N \mathbf{z}_n = \mathbf{0}$. The goal is to project the data onto an M -dimensional space ($M < D$) of M latent variables while maximizing the variance of the projected data.

To start, consider a one-dimensional projection ($M = 1$). We define the direction of this space using a D -dimensional unit³ vector $\mathbf{v}_1$, such that $\mathbf{v}_1^H \mathbf{v}_1 = 1$, where $\mathbf{v}_1^H$ denotes the conjugate transpose of $\mathbf{v}_1$. The data points $\mathbf{z}_n$ are then projected onto this direction, $\mathbf{v}_1^H \mathbf{z}_n$, and the variance of the projected data is

$$\frac{1}{N-1} \sum_{n=1}^{N-1} (\mathbf{v}_1^H \mathbf{z}_n - \mathbf{v}_1^H \bar{\mathbf{z}}) (\mathbf{v}_1^H \mathbf{z}_n - \mathbf{v}_1^H \bar{\mathbf{z}})^H = \mathbf{v}_1^H \mathbf{S}_{ZZ} \mathbf{v}_1, \quad (1.1)$$

where $\mathbf{S}_{ZZ}$ is the covariance matrix of $\mathbf{Z}$ given by

$$\mathbf{S}_{ZZ} = \frac{1}{N-1} \sum_{n=1}^N \mathbf{z}_n \mathbf{z}_n^H = \frac{1}{N-1} \mathbf{Z}^H \mathbf{Z}. \quad (1.2)$$

The goal is to find the vector $\mathbf{v}_1$ that maximizes this variance, formulated as:

$$\underset{\mathbf{v}_1}{\operatorname{argmax}} \quad \mathbf{v}_1^H \mathbf{S}_{ZZ} \mathbf{v}_1 \quad \text{subject to} \quad \mathbf{v}_1^H \mathbf{v}_1 = 1. \quad (1.3)$$

To solve this optimization problem, we maximize the projected variance $\mathbf{v}_1^H \mathbf{S}_{ZZ} \mathbf{v}_1$ with respect to $\mathbf{v}_1$. To enforce the normalization constraint $\mathbf{v}_1^H \mathbf{v}_1 = 1$, a Lagrange multiplier λ_1 is introduced, leading to the unconstrained maximization of:

³This normalization is for convenience and not necessary, since we are only interested in the direction of $\mathbf{v}_1$, not its magnitude.

$$\mathbf{v}_1^H \mathbf{S}_{ZZ} \mathbf{v}_1 + \lambda_1 (1 - \mathbf{v}_1^H \mathbf{v}_1). \quad (1.4) \quad 1$$

By setting the derivative with respect to $\mathbf{v}_1$ to zero, we find a stationary point when: 2

$$\mathbf{S}_{ZZ} \mathbf{v}_1 = \lambda_1 \mathbf{v}_1. \quad (1.5) \quad 3$$

This indicates that $\mathbf{v}_1$ must be a eigenvector of $\mathbf{S}_{ZZ}$, with λ_1 being the corresponding real-valued⁴ eigenvalue. The variance $\mathbf{v}_1^T \mathbf{S}_{ZZ} \mathbf{v}_1$ reaches its maximum when $\mathbf{v}_1$ aligns with the eigenvector associated with the largest eigenvalue⁵ λ_1 . The eigenvector is referred to as the first **vector of coefficients**⁶ (VC) of the data set. 4
5
6
7

To find additional VCs, we iteratively select new directions that maximize the projected variance while being orthogonal to the previously chosen directions. For an M-dimensional projection space, the optimal linear projection that maximizes the variance of the projected data is defined by the M eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_M$ of the covariance matrix $\mathbf{S}_{ZZ}$, corresponding to the M largest eigenvalues $\lambda_1, \dots, \lambda_M$. This result can be derived using proof by induction, leading to the concise expression: 8
9
10
11
12
13

$$\mathbf{S}_{ZZ} \mathbf{V} = \mathbf{V} \mathbf{\Lambda}, \quad (1.6) \quad 14$$

where $\mathbf{V}$ is unitary, i.e. $\mathbf{V}^T \mathbf{V} = \mathbf{I}_M$. 15

Thus, the eigenvectors of $\mathbf{S}_{ZZ}$ emerge as the new coordiante system "along" regularities in the data. For cases where these regularities are not aligned in unitary directions, see Section 1.1.2. 16
17

According to the spectral theorem for normal matrices [e.g. 51, p. 133], the hermitian covariance matrix $\mathbf{S}_{ZZ}$ can be decomposed into the sum of the outer products of the eigenvectors $\mathbf{v}_i$ and the corresponding real-valued eigenvalues λ_i : 18
19
20

⁴This follows from the property of the covariance matrix being hermitian.

⁵This can easily verified by left-multiplying Equation (1.5) with $\mathbf{v}_1^H$ and making use of the normalization condition.

⁶Naming conventions can vary significantly across different scientific domains. For example, in machine learning, the terms principal components and principal component scores refer to the eigenvectors (here $\mathbf{V}$) and their corresponding data projections (here $\mathbf{P}$) [e.g. 6]. In statistics, these are often called vectors of coefficients/loadings and principal components [e.g. 60], while in climate science and meteorology, they are known as empirical orthogonal functions and principal components/expansion coefficients [e.g. 101]. In this thesis, I aim to maintain consistency with the terminology commonly used in statistics.

$$\mathbf{S}_{ZZ} = \sum_{i=1}^M \lambda_i \mathbf{v}_i \mathbf{v}_i^H. \quad (1.7)$$

Each eigenvalue λ_i represents the amount of variance captured by the corresponding component, typically expressed as a fraction γ_i of the total variance:

$$\gamma_i = \frac{\lambda_i}{\sum_{i=1}^M \lambda_i} \quad (1.8)$$

Projecting the data $\mathbf{Z}$ onto the M -dimensional subspace thus provides a lower dimensional representation, called the principal components (PCs),

$$\mathbf{P} = \mathbf{Z}\mathbf{V}, \quad (1.9)$$

which also enables the reconstruction of the original data from the projected data:

$$\mathbf{Z} = \mathbf{P}\mathbf{V}^H. \quad (1.10)$$

However, calculating the eigenvector decomposition is not always computationally efficient. A more efficient method involves using Singular Value Domposition (SVD) [e.g. 51, p. 150], which decomposes the data matrix as

$$\mathbf{Z} = \mathbf{U}'\mathbf{\Sigma}\mathbf{V}'^H, \quad (1.11)$$

where $\mathbf{U}'$ and $\mathbf{V}'$ are unitary matrices, $\mathbf{U}'^H\mathbf{U}' = \mathbf{V}'^H\mathbf{V}' = \mathbf{I}_M$, and $\mathbf{\Sigma}$ is a diagonal matrix containing the singular values σ_i of $\mathbf{Z}$. The covariance matrix $\mathbf{S}_{ZZ}$ can then be expressed in terms of the SVD as

$$\mathbf{S} = \frac{1}{N-1} \mathbf{Z}^H \mathbf{Z} = \frac{1}{N-1} \mathbf{V}' \mathbf{\Sigma}^H \mathbf{U}'^H \mathbf{U}' \mathbf{\Sigma} \mathbf{V}'^H = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^H, \quad (1.12)$$

where, by comparison with (1.6), $\mathbf{V} = \mathbf{V}'$ and $\mathbf{\Lambda} = \frac{1}{N-1} \mathbf{\Sigma}^T \mathbf{\Sigma}$. This establishes a direct relationship between the eigenvalues of the covariance matrix and the singular values of the data matrix:

$$\lambda_i = \frac{\sigma_i^2}{N-1}. \quad (1.13) \quad 1$$

Comparing (1.10) with (1.11) shows that the PC scores $\mathbf{P}$ can be normalized 2

$$\mathbf{U} \equiv \mathbf{U}' = \mathbf{P}\mathbf{\Sigma}^{-1} = \frac{1}{(N-1)^{1/2}} \mathbf{P}\mathbf{\Lambda}^{-1/2}, \quad 3$$

$$\text{or alternatively } \mathbf{P} = (N-1)^{1/2} \mathbf{U}\mathbf{\Lambda}^{1/2}. \quad (1.14) \quad 4$$

Thus, $\mathbf{U}^H \mathbf{U} = \mathbf{I}_M$ ensures that the PC scores $\mathbf{U}$ are unitary and uncorrelated. 5

A common scaling method for the vector of coefficients $\mathbf{V}$ involves scaling by the square root of the corresponding eigenvalues, referred to as **loadings**: 6

$$\tilde{\mathbf{V}} = \mathbf{V}\mathbf{\Lambda}^{1/2}. \quad (1.15) \quad 8$$

Although not strictly unitary, these loadings form a generalized unitary matrix, as shown by 9

$$\tilde{\mathbf{V}}^H \tilde{\mathbf{V}} = \mathbf{\Lambda}. \quad (1.16) \quad 10$$

Using this scaling, the Equation 1.10 can be expressed as 11

$$\mathbf{X} = (N-1)^{1/2} \mathbf{U}\tilde{\mathbf{V}}^T. \quad (1.17) \quad 12$$

Computational Considerations The computational cost of performing eigendecomposition on the covariance matrix scales with $O(D^3)$, making it impractical for modern datasets, such as those in climate science, where typical data sizes are on the order of $D = O(10^6)$ or even larger. For example, ERA5 [16], a widely used reference for state-of-the-art reanalysis climate data, distributes a single climate variable on a regular 0.25° grid, resulting in $D = 1\,036\,800$ grid cells. 13

In practice, more efficient algorithms, such as SVD or power methods [39], have been employed traditionally. Recently, randomized low-rank approximation techniques [29, 42] have further reduced computational complexity to $O(ND \log(M) + (N+D)M^2)$. These advancements offer feasible alternatives for handling large, out-of-memory datasets (see Table 1.1). 14
15
16
17
18
19
20
21

Shortcomings Despite its widespread use and versatility, PCA has several limitations and potential drawbacks that should be considered:

1. **Linearity Assumption:** PCA assumes a linear structure in the data, which may not hold in complex systems such as climate data, where nonlinear relationships can be significant [64, 99].
2. **Interpretability of VCs:** The orthogonality of the vectors of coefficients is a mathematical construct that may not always correspond to the physical reality of the data, potentially leading to misinterpretations [94].
3. **Variance Focus:** Focusing on directions of maximum variance may not always capture the most relevant features for a specific application, potentially critical but less variable information [35].
4. **Static Nature:** As a static method, PCA does not account for temporal dynamics or changes over time, which is particularly limiting in climate science, where time evolution is a key aspect.
5. **Sensitivity to Outliers:** PCA can be sensitive to outliers, as extreme values can disproportionately influence the VCs, leading to misleading results [19].
6. **Sensitivity to Scaling:** The results are sensitive to the scaling of the input data, meaning that variables with larger ranges can dominate the VCs unless the data is appropriately standardized.
7. **Loss of Information:** While dimensionality reduction is often the goal, PCA inherently discards some information, which could lead to the loss of important details crucial for accurate analysis and interpretation in specific applications.

Over the past few decades, numerous modifications and extensions have been proposed to address these limitations, especially in the context of climate science. For instance, to address the linearity assumption, various forms of kernel PCA [13, 99, 112] have been developed. For improved interpretability, rotated [94] and sparse [61] PCA methods have been proposed. Robust PCA [19] offers a more resilient approach to handling sensitivity to outlier. Furthermore, extended [11, 111] and complex [3, 49] PCA methods have been introduced to account for temporal dynamics. In addition, eigenmethods like Optimal Persistence Analysis (OPA) [26, 27] and Principal Oscillation Pattern (POP) analysis [46] provide alternative approaches independent to variance maximization to understanding dominant patterns in physical climate fields.

In the following sections, we will explore some of these shortcomings and proposed solutions in greater detail. Specifically, we will discuss how additional rotation of the PCA solution can

Table 1.1 — Computational costs of different algorithms used to compute either full or approximative M -rank solutions for PCA for a real, dense, in-memory matrix with dimensions $N \times D$. Compiled from Bishop [6] and Halko *et al.* [42].

Algorithm	Complexity
Eigenvector decomposition	$O(D^3)$
Full SVD	$O(ND \min(N, D))$
Power method	$O(MD^2)$
Truncated SVD	$O(NDM)$
Randomized SVD	$O(ND \log(M) + (N + D)M^2)$

enhance physical interpretability (Section 1.1.2) and how complexifying the data can provide insights into time evolution patterns (Section 1.1.3).

Weighting of Data Variables In some cases, it may be necessary to weight the data variables differently. This can be achieved by scaling the data matrix $\mathbf{Z}$ by a diagonal matrix $\mathbf{W}$, where the j^{th} diagonal element w_j corresponds to the weight of the j^{th} variable. A very common situation is that the observed variables exhibit stark differences in their variance, which can sometimes lead to mask structure in the data as the PCA solution is solely driven by the different scales of variance (so the first principal axis will closely align to the variable with the highest variance) in which case the obvious choice is to use the inverse of the standard deviation of each variable, $w_j = 1/\sigma_j$, which ensures that all variables are equally weighted in which case the covariance matrix becomes the correlation matrix. Another common situation in climate science is that the physical fields are represented on regular longitude-latitude grids. This representation leads to a higher density of grid points near the poles, which typically results in an exaggerated importance of components spatially focused near the poles. To address this, the data matrix can be weighted by the square root of the cosine of the latitude [82], $w_j = \cos(\text{lat}_j)$, where lat_j is the latitude of the j^{th} grid point. This weighting ensures that the grid points are equally weighted in the PCA solution, regardless of their geographical location. In the following we will discuss some popular extensions of PCA together with a related multi-set method called Maximum Covariance Analysis (see Section 1.2), for which the latitude weighting can be applied in the same way.

1.1.2 Sparse PCA via Rotation

PCA is widely used for dimensionality reduction, as it maximizes the explained variance and provides orthogonality, making it optimal in a least-squares sense [6]. It effectively identifies principal axes that capture the covariance among variables, allowing for the examination of interrelationships between stations, grid points or variables. However, PCA is not ideal for identifying the cause of variation or deriving physical modes of variability [60, 69, 94].

One limitation is that real-world processes do not necessarily adhere to the orthogonality imposed by PCA's eigenvectors. Additionally, the dense nature of these VCs – where all elements are non-zero – can make interpretation challenging. In the context of climate and environmental science, where variables are often represented across spatial grids, these limitations can lead to recognizable but artificial patterns known as Buell patterns [12, 68, 94]. These patterns typically start with a positive component and evolve into increasingly complex structures, such as bipole or tripole patterns (Figure 1.1). Such issues highlight the disconnect between standard PCA and underlying physical processes [94], prompting the development of alternative methods that produce sparse VCs, characterized by many zero or near-zero elements. This sparsity – which necessarily implies the loss of orthogonality for at least one of the VCs and PC scores – can enhance interpretability by localizing the patterns and reducing the density of the VCs.

In climate science, the most common approach to achieving sparse PCA is through rotation techniques [50, 93, 94], a concept borrowed from factor analysis [cf. 60, pp. 269]. In this "factor analysis viewpoint" of PCA [60, pp. 150-161], a linear transformation is applied to a subset of the VCs to derive a new set of rotated VCs that are "simpler" in a heuristically defined manner, ideally aligning better with the underlying physical processes.

As noted in Equation (1.17), once the orthogonality constraint is relaxed, PCA solutions exhibit inherent indeterminacy, allowing any non-singular transformation $\mathbf{R}$ to satisfy the PCA formulation [e.g. 94, p. 305]:

$$\mathbf{Z} = (\mathbf{N} - 1)^{1/2} \mathbf{U} \tilde{\mathbf{V}}^H = (\mathbf{N} - 1)^{1/2} \mathbf{U} \mathbf{R} \mathbf{R}^{-1} \tilde{\mathbf{V}}^H. \quad (1.18)$$

Rotated PCA exploits this indeterminacy by choosing $\mathbf{R}$ to optimize a simplicity criterion f , resulting in rotated loadings⁷ $\tilde{\mathbf{V}}_R = \tilde{\mathbf{V}} \mathbf{R}$ that maximize f :

$$\underset{\mathbf{R}}{\operatorname{argmax}} f(\tilde{\mathbf{V}}_R). \quad (1.19)$$

⁷In principle, it is possible to rotate the unitary matrix of coefficients $\mathbf{V}$ instead of the loadings matrix $\tilde{\mathbf{V}}$ [60, 77]. However, this approach is seldom used in practice. The primary reason is that rotating the unitary coefficient matrix preserves the orthogonality of the VCs but results in the loss of orthogonality among the PCs, which does not address issues like the Buell patterns. Moreover, this method does not consider the relative importance of the spatial patterns, as indicated by the eigenvalues, making the results of the rotation process highly sensitive to the selection of the number of components to rotate.

There are two main categories of rotation: *orthogonal*⁸ and *oblique*. In orthogonal rotation, $\mathbf{R}$ is a unitary matrix, preserving the uncorrelated nature of the PCs since $\mathbf{U}_R^H \mathbf{U}_R = \mathbf{R}^H \mathbf{U}^H \mathbf{U} \mathbf{R} = \mathbf{I}_M$. In oblique rotation, $\mathbf{R}$ is a general non-singular matrix, resulting in correlated PCs, as $\mathbf{U}_R^H \mathbf{U}_R = \mathbf{R}^H \mathbf{R}$. Regardless of the method, the rotated loadings are not orthogonal,

$$\tilde{\mathbf{V}}_R^H \tilde{\mathbf{V}}_R = \mathbf{R}^H \tilde{\mathbf{V}}^H \tilde{\mathbf{V}} \mathbf{R} \stackrel{(1.16)}{=} \mathbf{R}^H \mathbf{A} \mathbf{R} \neq \mathbf{I}_M, \quad (1.20)$$

which is crucial for achieving a more interpretable solution. Various simplicity criteria f have been proposed [see e.g. 94], with most focusing on making the rotated loadings more localized and physically interpretable. In practice, the specific choice of simplicity criterion often has little impact on the results [94].

Varimax Orthogonal Rotation The most widely used orthogonal rotation technique is the varimax criterion [62], which aims to simplify the PCs by maximizing the variance of the squared loadings,

$$f(\tilde{\mathbf{V}}_R) = \sum_{m=1}^M \left(p \sum_{j=1}^P b_{jm}^2 - \left(\sum_{j=1}^P b_{jm} \right)^2 \right) \quad (1.21)$$

where b_{jm} are the squared norm of the rotated loadings $|\tilde{v}_{R,jm}|^2/h_j^2$ normalized by the communalities $h_j^2 = \sum_{m=1}^M |\tilde{v}_{R,jm}|^2$ of variable j . The varimax rotation pushes the elements of $\tilde{\mathbf{V}}_R$ towards either 0 or 1, creating a sparse solution with primarily large or small values, which is particularly useful for identifying localized physical processes, such as atmospheric or oceanic teleconnections.

Promax Oblique Rotation The promax criterion [48], a generalization of the varimax criterion, is one of the most common oblique rotation techniques. It is based on the premise that an orthogonal simple structure solution, such as varimax, is often close to an optimal oblique solution. The promax method constructs a hypothetical oblique target matrix $\mathbf{T}$ by performing an oblique Procrustes transformation [57] on the varimax-rotated loadings matrix,

⁸In the broader context of complex data, this would be referred to as a unitary rotation. However, since factor analysis has primarily been applied to real-valued data, the terms orthogonal and oblique rotations have become the standard. To avoid confusion, I use these historical terms while acknowledging that they also encompass the more general case of complex data.

$$\mathbf{T} = \tilde{\mathbf{V}}_{\mathbf{R}} \mathbf{R} + \mathbf{E}, \quad (1.22)$$

where $\mathbf{R}$ is the transformation matrix and $\mathbf{E}$ is a matrix of residuals. The normal equation for solving the complex Procrustes problem is given by:

$$\mathbf{R} = (\tilde{\mathbf{V}}_{\mathbf{R}}^H \tilde{\mathbf{V}}_{\mathbf{R}})^{-1} \mathbf{T}^H \tilde{\mathbf{V}}_{\mathbf{R}}. \quad (1.23)$$

To define the target matrix $\mathbf{T}$, the promax criterion increases the discrimination between high and low loadings by raising the loadings to a power q while retaining the original sign,

$$t_{jm} = \frac{|a_{jm}|^{q+1}}{a_{jm}}, \quad (1.24)$$

where q controls the degree of obliqueness and a_{jm} are the elements of the max-normalized, varimax-rotated loadings matrix $\mathbf{A}$, defined as

$$\mathbf{A} = \tilde{\mathbf{A}} \mathbf{K}^{-1}, \quad (1.25)$$

$$\tilde{\mathbf{A}} = \mathbf{H}^{-1} \tilde{\mathbf{V}}_{\mathbf{R}}, \quad (1.26)$$

with $\mathbf{K}$ being a diagonal matrix containing the maximum absolute values of each column of $\tilde{\mathbf{A}}$, $a_m = \max_m |A_{pm}|$, and $\mathbf{H}$ a diagonal matrix with the square root of the communalities h_j for the j^{th} variable.

Application and Criticism Both orthogonal and oblique rotation techniques, particularly varimax and promax, have been extensively used since the 1980s to identify stationary modes of variability in climate data [e.g. 4, 38, 50, 58, 67, 72, 78]. Notable climate oscillations, such as the North Atlantic Oscillation (NAO), have been defined based on varimax-rotated patterns of loadings [4]. These techniques also help identify "ghost patterns", artificial patterns identified by PCA that are not physically meaningful [58, 71].

Rotated modes offer several advantages, such as greater robustness against temporal and spatial subsampling [21] compared to unrotated modes. They also effectively mitigate the statistical artifacts associated with Buell patterns, resulting in more realistic representations of the underlying autocorrelation structure and thus increased interpretability [71, 94, 95].

Moreover, rotation techniques have been successfully interated into other PCa extensions dealing with the non-linear and non-stationary nature of climate oscillations [13, 41, 49].

However, rotation techniques are not without criticism, mainly due to the ad-hoc nature of the various simplicity criteria [cf. 60, pp. 270-274]. This has led to the exploration of alternative approaches to enforce sparsity. Notably, Jolliffe *et al.* [61] introduced the Simplified Component Technique - LASSO (SCoTLASS), a sparse PCA variant based on the least absolute shrinkage and selection operator (LASSO), originally from regularized regression [104]. This method imposes a constraint on the ℓ_1 norm of the vector of loadings, resulting in some loading coefficients being exactly zero.

Despite its advantages, SCoTLASS has drawbacks, particularly when the number of variables D is much larger than the number of observations N . In such cases, SCoTLASS can only select at most N non-zero elements in the loading vectors, which is not ideal in climate science where often $D \gg N$. Additionally, since SCoTLASS solutions cannot be obtained via SVD and require a numerical approach using projected gradient methods [106], the computational demand is typically higher than for (rotated) PCA. While this method has inspired many other sparse PCA techniques based on penalties or constraints (usually ℓ_0 or ℓ_1) [see e.g. 34, 115, 116], these formulations are often complex and require algorithmic innovations to avoid issues like local optima and ensure convergence. They may also face challenges with memory limitations, slow computation speeds and limited effectiveness in modelling the data [18, 40].

It is likely for these reasons, that rotated PCA remains one of the most commonly used techniques for achieving sparse PCA solutions in climate science. It is computationally efficient, straightforward to implement and provides a more interpretable understanding of the data compared to standard PCA.

1.1.3 Hilbert PCA

While PCA efficiently identifies the principal axes along which the covariance among all variables is distributed, and rotated PCA provides a more sparse solution related to underlying physical processes, both methods are inherently static. They do not for temporal dynamics or changes over time, as they only consider the current state of the physical field. This limitation is significant in climate science, where time evolution is crucial. For instance, standard PCA does not differentiate between the order of data points over time, resulting in the same eigenvectors and eigenvalues regardless of temporal permutations.

This static nature poses a challenge for capturing oscillatory patterns, such as QBO, MJO, ENSO and BSISO, which involve spatial patterns propagating over time. These phenomena cannot be adequately captured by standard PCA.




Figure 1.1 — Comparison of the first three modes obtained from PCA vs. Varimax-rotated PCA for a synthetic data set. (Left) Original data set with three different, uncorrelated autoregressive processes. (Middle) PCA solution. (Right) Varimax-rotated PCA solution.

1 Hilbert Transform as a Momentum Operator

To address this limitation, one approach is to model both the current state of the climate and its tendency or ‘momentum’, analogous to the Hamiltonian principle in mechanics, where a system’s future is described by first-order differential equations for location (state) and momentum (tendency) [cf. 101, p. 353]. In classical mechanics, propagating features are represented using canonical coordinates: each state variable is augmented by a momentum variable, allowing oscillations to be depicted as rotations in a two-dimensional phase space. From a data perspective, this augmentation can be achieved via the Hilbert transform ($\mathcal{H}$). For a real time series $x(t)$, the Hilbert transform $x'(t)$ is defined as

$$x'(t) \equiv \mathcal{H}\{x\}(t) = \frac{1}{\pi} \text{p.v.} \int_{-\infty}^{\infty} \frac{x(\tau)}{\tau - t} d\tau, \quad (1.27)$$

where p.v. denotes the Cauchy principal value ensuring convergence despite the singularity at $\tau = t$. The Hilbert transform emphasises the local properties of $x(t)$ by convolution with $\frac{1}{\pi t}$. In the frequency domain, it is related to the Fourier transform $\mathcal{F}$ by

$$\mathcal{F}\{y\}(\omega) = -i \operatorname{sgn}(\omega) \mathcal{F}\{x\}(\omega) \quad (1.28)$$

where the imaginary unit i indicates a phase shift of $\pi/2$ for the frequency components ω of x , placing x' in quadrature to x . This provides an intuitive interpretation of the Hilbert transform, as it acts as a linear filter that alters the phase of the signal's components by 90 degrees without changing their amplitude.

By using the Hilbert transform, a real-valued time series $x(t)$ can be complexified into $z(t) = x(t) + i\mathcal{H}\{x\}(t)$, known as the *analytic signal*, which allows for the computation of the instantaneous amplitude and phase of $x(t)$:

$$A(t) = |z(t)| = \sqrt{x(t)^2 + \mathcal{H}\{x\}(t)^2}, \quad \phi(t) = \arg(z(t)) = \arctan\left(\frac{\mathcal{H}\{x\}(t)}{x(t)}\right). \quad (1.29)$$

Similar results can be obtained for the discrete case [84, pp. 942-956], and this concept can be extended to multivariate data [e.g. 43, pp. 105-106], defining the analytic signal of a data matrix X as

$$Z = X + i\mathcal{H}\{X\}, \quad (1.30)$$

where $\mathcal{H}\{X\}$ is the feature-wise Hilbert transform of X .

Rasmusson *et al.* [91] introduced the concept of Hilbert PCA using this data augmentation technique⁹. This method, refined by Barnett [3] and Horel [49], involves applying PCA to the augmented data matrix Z . By incorporating phase information, Hilbert PCA can capture the direction and speed of wave-like patterns, unlike traditional PCA, which only identifies standing oscillations. The resulting eigenvectors and principal component scores are complex, with the real and imaginary components linked via the Hilbert transform. The variance capture by each component is represented by the real eigenvalues of the complex covariance matrix of Z (see Section 1.1).

The main advantage of Hilbert PCA is the ability to identify and reconstruct propagating waves, which is especially valuable in studying climate phenomena characterized by such patterns, including MJO, ENSO, QBO and BSISO. For instance, Cai & Baines [17] utilized Hilbert PCA to analyze SST in the southern hemisphere, linking the findings to ENSO, while Bürger [15] employed it to develop an index for BSISO. However, to reliably identify oscillatory patterns, sufficiently long time series are required to resolve at least one full oscillation [76]. Additionally, interpreting of the Hilbert transform as representing the instantaneous frequency is valid holds for narrow-band signals; the phase of the Hilbert transform is not well-defined

⁹Although Rasmusson *et al.* [91] originally termed it *Hilbert Singular Decomposition*, the terms *Hilbert EOF/PCA* and *Complex EOF analysis* are more commonly used in the literature [15, 43, 60, 101].

for broadband signals [8, 37]. Moreover, real-time monitoring is challenging with Hilbert PCA because new real-valued data cannot be directly projected onto the existing patterns without knowing or estimating the exact momentum (phase) of the new data, which requires future information (Equation (1.27)). To mitigate this, an approximative real-time (causal) Hilbert transform can be employed [14, 15], offering a practical remedy to the issue of real-time applicability. Finally, Hilbert PCA shares PCA's limitations regarding physical interpretability due to the orthogonality of the VCs. To address this, rotated Hilbert PCA has been proposed to enhance interpretability [7, 50].

1.2 Maximum Covariance Analysis

In contrast to finding latent variables that maximize covariance within a single dataset, Maximum Covariance Analysis (MCA) aims to identify the latent variables that are common in two sets of data and maximize covariance. For example, in the case of the MJO [75], which is the dominant mode of intraseasonal variability in the tropical atmosphere, characterized by an eastward-moving cell of overturning wind circulation and associated regions of enhanced and suppressed precipitation, MCA helps in findings shared dynamics between different variables. This method allows us to focus on shared variability while ignoring dominant variations associated with only one variable, such as diurnal cycles linked to surface processes absent at higher atmosphere levels.

MCA identifies patterns of maximum covariance between datasets by finding pairs of linear subspaces with high cross-covariance [e.g. 101, p. 321]. Unlike PCA, which operates on a single set of variables, MCA considers two¹⁰ sets of variables and seeks components that exhibit high covariance across datasets. This approach can be formulated as an optimization problem, aiming to find pairs of (complex-valued) coupled patterns $\mathbf{v}_{s,1}$ for datasets $\mathbf{Z}_s$, $s = x, y$, such that the projection maximizes the covariance between the datasets. The solution involves a generalized eigenvector problem, often solved using the Singular Value Decomposition (SVD) of the cross-covariance matrix.

$$\operatorname{argmax}_{\mathbf{v}_{x,1}, \mathbf{v}_{y,1}} \mathbf{v}_{x,1}^T \mathbf{C}_{xy} \mathbf{v}_{y,1} \quad \text{subject to} \quad \|\mathbf{v}_{x,1}\|_2 = \|\mathbf{v}_{y,1}\|_2 = 1, \quad (1.31)$$

where the cross-covariance matrix between the two datasets is given by

¹⁰In theory, this can be extended to more than two sets of data [63].

$$\mathbf{C}_{xy} = \frac{1}{N-1} \mathbf{Z}_x^H \mathbf{Z}_y. \quad (1.32) \quad 1$$

The solution to this problem gives rise to a generalized eigenvalue equations similar to that of PCA, which can also be solved by using the SVD of the cross-covariance matrix 2
3

$$\mathbf{C}_{xy} = \mathbf{V}_x \mathbf{\Sigma} \mathbf{V}_y^H \quad (1.33) \quad 4$$

where the i^{th} column of the left and right singular vectors, $\mathbf{V}_x$ and $\mathbf{V}_y$, contains the pattern vectors $\mathbf{v}_{x,i}$ and $\mathbf{v}_{y,i}$, and $\mathbf{\Sigma}$ a diagonal matrix containing the singular values σ_i in decending order. The first left and right singular vectors correspond to the maximum covariance, given by the largest singular value σ_1 . 5
6
7
8

Since the pattern vectors are unitary, datasets can be expanded as 9

$$\mathbf{Z}_s = \mathbf{P}_s \mathbf{V}_s^H, \quad (1.34) \quad 10$$

with the expansion coefficients $\mathbf{P}_s$. The expansion coefficients $\mathbf{P}_x$ and $\mathbf{P}_y$ are uncorrelated with each other: 11
12

$$\mathbf{P}_x^H \mathbf{P}_y = \mathbf{V}_x^H \mathbf{Z}_x^H \mathbf{Z}_y \mathbf{V}_y \stackrel{(1.32)}{=} (N-1) \mathbf{V}_x^H \mathbf{C}_{xy} \mathbf{V}_y \stackrel{(1.33)}{=} (N-1) \mathbf{\Sigma}. \quad (1.35) \quad 13$$

However, the expansion coefficients $\mathbf{P}_x$ and $\mathbf{P}_y$ are generally correlated within themselves as indicated by the condition $\mathbf{P}_s^H \mathbf{P}_s \neq c \mathbf{I}_M$ for any arbitrary constant c . 14
15

The total covariance, given by the squared Frobenius norm of the cross-covariance matrix, is 16

$$\|\mathbf{C}_{xy}\|_F^2 = \sum_{i=1}^M \sigma_i^2, \quad (1.36) \quad 17$$

which shows that each mode captures a fraction of the total squared covariance. The relative importance of each mode can be quantified by the squared covariance fraction $f_{xy,i}$: 18
19

$$f_{xy,i} = \frac{\sigma_i^2}{\sum_{i=1}^M \sigma_i^2}. \quad (1.37)$$

This fraction provides a measure of how much of the shared variability between the two datasets is captured by the i^{th} mode. To assess the importance of the patterns themselves, it is necessary to consider the variance within the individual datasets explained by each mode relative to the total variance of the individual datasets. This can be expressed using the squared Frobenius norm of the individual datasets:

$$f_{xx,i} = \frac{\mathbf{v}_{x,i}^H \mathbf{C}_{xx} \mathbf{v}_{x,i}}{\|\mathbf{Z}_x\|_F^2} = \frac{1}{N-1} \frac{\mathbf{p}_{x,i}^H \mathbf{p}_{x,i}}{\|\mathbf{Z}_x\|_F^2}, \quad (1.38)$$

$$f_{yy,i} = \frac{\mathbf{v}_{y,i}^H \mathbf{C}_{yy} \mathbf{v}_{y,i}}{\|\mathbf{Z}_y\|_F^2} = \frac{1}{N-1} \frac{\mathbf{p}_{y,i}^H \mathbf{p}_{y,i}}{\|\mathbf{Z}_y\|_F^2}. \quad (1.39)$$

While the squared covariance fraction $f_{xy,i}$ indicates the amount of shared variability between the two datasets, it does not provide information about the strength of the association or coupling between them. To measure this, we use the correlation coefficient ρ_i for the expansion coefficients of mode i between the two datasets:

$$\rho_i = \frac{\text{cov}(\mathbf{p}_{1,i}, \mathbf{p}_{2,i})}{(\text{var}(\mathbf{p}_{1,i}) \text{var}(\mathbf{p}_{2,i}))^{1/2}} = \frac{\sigma_i}{\|\mathbf{p}_{1,i}\|^2 \|\mathbf{p}_{2,i}\|^2}. \quad (1.40)$$

The homogeneous patterns [10] for the pattern vectors are:

$$\text{HoP}_{x,i} = \text{cov}(\mathbf{Z}_x, \mathbf{p}_{x,i}) = \mathbf{S}_{xx} \mathbf{v}_{x,i}, \quad (1.41)$$

$$\text{HoP}_{y,i} = \text{cov}(\mathbf{Z}_y, \mathbf{p}_{y,i}) = \mathbf{S}_{yy} \mathbf{v}_{y,i}, \quad (1.42)$$

These patterns provide an indicator of which variables or features are important for the covarying components in both fields. The corresponding heterogeneous patterns are defined as:

$$\text{HeP}_{x,i} = \text{cov}(\mathbf{Z}_x, \mathbf{p}_{y,i}) = \sigma_i \mathbf{v}_{y,i}, \quad (1.43)$$

$$\text{HeP}_{y,i} = \text{cov}(\mathbf{Z}_y, \mathbf{p}_{x,i}) = \sigma_i \mathbf{v}_{x,i}. \quad (1.44)$$

These patterns indicate how well the variables in one field can be predicted from the i^{th} expansion coefficient of the other field. One may also use standardized data, $\mathbf{Z}_x$ and $\mathbf{Z}_y$, to obtain the respective homogeneous and heterogeneous correlation patterns.

MCA gained popularity in the field of meteorology¹¹ through the works of Bretherton *et al.* [10] and Wallace *et al.* [109], since when it has been extensively applied [e.g. 5, 70, 81, 92, 107, 109].

The method's orthogonal patterns ensure efficient dimensionality reduction of the covariance but limit physical interpretability, similar to PCA. To enhance interpretability, rotation of the MCA solution can be employed, as proposed by Cheng & Dunkerton [20], and later applied by Frankignoul & Kestenare [36]. A composite matrix of the left and right singular vectors is constructed and a simplicity criterion (e.g. varimax) is applied:

$$\mathbf{V}_R = \begin{bmatrix} \mathbf{V}_x \\ \mathbf{V}_y \end{bmatrix} \mathbf{M} \quad (1.45)$$

In this case, $\mathbf{M}$ is a diagonal matrix that weights the augmented pattern matrix. While Cheng & Dunkerton [20] proposed $\mathbf{M} = \Sigma^{1/2}$, this choice is not ideal as the singular values σ_i are not directly related to the explained covariance of the patterns. Instead, it is more appropriate to weight each mode by the squared singular value, preserving the squared covariance fraction of the first m modes before and after (orthogonal) rotation. Thus, we use

$$\mathbf{M} = \Sigma \quad (1.46)$$

for the weighted (or "loaded") pattern vectors.

Beyond rotation for better interpretability, Hilbert MCA can be applied to capture the temporal evolution of coupled patterns, akin to Hilbert PCA (see Section 1.1.3). This is especially useful for analyzing propagating patterns like the MJO, which involve covarying patterns in

¹¹According to Jolliffe [60, p. 225], the technique was first used in psychology in the 1950s under the name *inter-battery factor analysis*. As a variant of Partial Least Squares (PLS), it was later applied in biomedical sciences by Sampson *et al.* [97] and was referred to as *PLS-SVD* by Wegelin [113]. In meteorology, Bretherton *et al.* [10] introduced the term *SVD analysis*, which remains in use today [e.g. 81]. Occasionally, the method is also called *Canonical Covariance Analysis*, similar to Canonical Correlation Analysis [e.g. 22, 43].

different climate variables, such as wind speed in the troposphere and lower stratosphere, and precipitation. By augmenting data with the Hilbert transform, we obtain complexified fields and solve the maximization problem (Equation (1.31)) for these augmented data, capturing the temporal evolution of coupled patterns. It seems that Hilbert MCA has seen limited application in climate science [e.g. 32]. In this thesis we utilize the combined power of Hilbert and rotated MCA to identify and interpret coupled patterns of variability in the climate system.

PCA, MCA and CCA MCA reduces to PCA when both sets are identical. Furthermore, MCA can be viewed as two PCAs followed by rotation to align the results of the two analyses as closely as possible [23]. MCA is also closely related to Canonical Correlation Analysis (CCA), which, instead of maximizing shared covariance, seeks patterns that maximize correlation. Unlike MCA, where pattern vectors are by construction orthogonal, yielding correlated expansion coefficients, CCA results in pattern vectors that are not necessarily orthogonal but produce uncorrelated canonical variates [60, 114]. In fact, PCA, MCA and CCA can all be seen as special cases of a more general framework [9, 102].

1.3 Neural Networks

1.3.1 Introduction to Neural Networks

Definition and Overview: Explain what neural networks are, including a brief history and their importance in machine learning. **Basic Architecture:** Describe the basic structure of a neural network, including neurons, layers (input, hidden, output), activation functions, and the learning process.

1.3.2 Recurrent Neural Networks

Introduction to RNNs¹²: Describe the architecture of RNNs, their ability to handle sequential data, and their applications. **Challenges with RNNs:** Explain issues such as vanishing and exploding gradients.

Long Short-Term Memory (LSTM) Networks

Introduction to LSTM¹³: Discuss the motivation behind LSTM networks and how they address the limitations of standard RNNs. **LSTM Architecture:** Describe the components of LSTM cells (input gate, forget gate, output gate) and their functions. **Mathematical Formulation:** Provide

¹²Graves 2021 Supervised Sequence Labelling with Recurrent Neural Networks

¹³Hochreiter & Schmidhuber 1997

the equations governing LSTM networks. Applications of LSTM: Highlight key applications where LSTMs have shown significant performance improvements.

1.3.3 Explainable AI Techniques

Introduction to Explainable AI: Explain the need for explainable AI, particularly in critical applications where understanding model decisions is crucial. Overview of Techniques: Briefly list various explainable AI methods, with a focus on LRP and SHAP.

Layerwise Relevance Propagation (LRP)

Concept of LRP¹⁴: Describe what LRP is and its importance in understanding neural network decisions. Mechanism of LRP: Explain how LRP works, including the propagation rules and the relevance score computation. Applications of LRP: Provide examples of how LRP is used in practice to interpret model predictions.

simplest continuous latent variable model assumes Gaussian distributions for both the latent and observed variables. simplest continuous latent variable model assumes Gaussian distributions for both the latent and observed variables

SHapley Additive exPlanations (SHAP)

Introduction to SHAP¹⁵: Discuss the SHAP method, its foundation in cooperative game theory, and its advantages. SHAP Values: Explain how SHAP values are computed and their interpretation. Applications of SHAP: Provide examples of SHAP in various domains, highlighting its ability to provide consistent and locally accurate explanations.

1.4 Gaussian Processes

Class of pattern recognition techniques that typically require a metric to be defined that measures the similarity of any two vectors in input space, and are generally fast to "train" but slow at making predictions for test data points.

¹⁴Bach et al. 2015

¹⁵Lundberg 2017

1.4.1 Introduction to Gaussian Processes

Definition and Overview: Explain what Gaussian Processes are, their significance in machine learning, and their use in modeling continuous data. Basic Concepts: Introduce key concepts such as Gaussian distributions, covariance functions, and kernels.

1.4.2 Gaussian Processes for Regression

GP Regression Model: Describe the formulation of GP for regression tasks, including prior, likelihood, and posterior distributions. Kernel Functions: Discuss common kernel functions (e.g., squared exponential, Matérn) and their role in defining the GP's properties. Hyperparameter Tuning: Explain how to tune hyperparameters using techniques like maximum likelihood estimation.

1.4.3 Spatial Modeling with Gaussian Processes

Spatial Data: Define spatial data and the importance of modeling spatial dependencies. GPs in Spatial Modeling: Explain how GPs are particularly suited for spatial data due to their ability to model complex, non-linear relationships and provide uncertainty estimates.

Log Gaussian Cox Process

Introduction to Cox Processes: Explain Cox processes and their relevance in spatial point process modeling. Log Gaussian Cox Model: Describe the log Gaussian Cox model, which assumes the log intensity function is modeled by a GP. Mathematical Formulation: Provide the mathematical formulation of the log Gaussian Cox model.

Negative Binomial Distribution

Introduction to Negative Binomial Distribution: Explain the negative binomial distribution and its use in modeling overdispersed count data. Integration with Log Gaussian Cox Model: Describe how the negative binomial distribution can be integrated into the log Gaussian Cox model to handle count data with overdispersion.

1.4.4 Model Implementation and Inference

Model Fitting: Discuss methods for fitting the log Gaussian Cox model with a negative binomial distribution, including Markov Chain Monte Carlo (MCMC) and variational inference.

1.5 Comparative Perspective

21

Inference and Prediction: Explain how to perform inference and make predictions using the fitted model.

1
2

1.5 Comparative Perspective

3

Bridging Theory and Scientific Practice in a Data-Driven World

2

Over the last decades, the field of climate science has seen a rapid increase in the volume and complexity of available data [85]. This trend is driven by the growing number of Earth observation satellites, the development of high-resolution climate models, and the increasing availability of citizen science data. As a result, researchers are faced with the challenge of analyzing vast amounts of data to extract meaningful insights about the Earth's climate and environment system. To address this challenge, dimensionality reduction methods have become an essential tool for gaining a better understanding of complex climate processes.

In the last chapter we briefly reviewed the theory for some commonly applied dimensionality reduction techniques in climate and environmental science. We showed that the combination of some useful and known extensions for PCA can be used to extend MCA to extract physically more meaningful patterns to analyze the shared dynamics between two distinct datasets. In this chapter, we will briefly describe

how the ever-growing volume of climate data produces practical challenges for researchers in the application of these techniques and how the open-source software ecosystem has responded to these challenges. Specifically, Section 2.3 will introduce the `xeofs` library, a Python package that aims to facilitate the application of scalable dimensionality reduction techniques in climate science. By providing a user-friendly interface to traditional and state-of-the-art methods specific to the climate science context, `xeofs` enables researchers to leverage the power of advanced data analysis techniques without the need for specialized expertise. This section is based upon the work presented in:

- **Rieger, N.** & Levang, S. J. `xeofs`: Comprehensive EOF Analysis in Python with Xarray. *Journal of Open Source Software* **9**, 6060. doi:[10.21105/joss.06060](https://doi.org/10.21105/joss.06060) (2024)

2.1 Scalable Open-Source Tools for Climate Data Analysis

Climate science is a data-intensive field that relies on the analysis of large and complex datasets to understand the Earth's climate system. The increasing volume of available data poses a




Figure 2.1 — Open-source software (OSS) projects for environmental science. Panel A shows the number of new OSS projects created per year with the relative importance of Python based projects. Panel B shows the relative share of programming languages used in all created OSS projects since 2010. Programming languages with a share $\leq 0.5\%$ have been excluded. OSS project data is made available from *Open Sustainable Technology* [1]. Here we defined environmental science, we considered all projects in the categories "Atmosphere", "Climate Change", "Cryosphere", "Hydrosphere" or "Biosphere". Note that for 2024 the data includes only the first half of the year until August.

significant challenge for researchers, as traditional data analysis methods are often ill-suited to handle the scale and complexity of climate data. Additionally, the application of some of these techniques requires specialized knowledge and expertise, making them inaccessible to many researchers in the field. To bridge the gap between theory and scientific practice, software tools that lower the barrier to entry for non-experts are essential. Consequently, the field of climate science has seen a growing interest in the development of open-source tools that enable scalable and efficient data analysis.

One of the biggest challenges in climate science today is the exploding volume of climate data and the resulting challenge of data analytics, i.e. the ability to extract meaningful information from large, complex datasets [98]. quotation: "the capacity to perform useful scientific analyses over large quantities of data in reasonable amounts of time." (Schnase 2016). In many respects this is the biggest challenge, for without effective means for transforming large scientific data collections into meaningful scientific knowledge, our climate science mission fails.

To avoid that the analysis of these large data sets remains in the hand of a selective group of "Wissende", it effectively requires free, robust and scalable software tools that allow researchers without specialized knowledge in programming to simplifying and managing the vast amounts of data produced by climate models [30].

At the same time, over the last decade the open-source software (OSS) movement has gained significant momentum in the scientific community (Fig. 2.1), with many researchers recognizing the benefits of OSS for fostering collaboration, improved transparency and reproducibility in scientific research [59, 65, 80]. In particular, the high-level, general-purpose programming language Python has emerged as one of the most popular tools for scientific computing (Fig. 2.1 B), due to its ease of use, extensive library support, and vibrant community of developers. Prime examples of the success of collaborative open-source software development in the scientific community are the *NumPy* [44, 83], *SciPy* [83, 108] and *Matplotlib* [55] libraries, which have become essential tools for scientific computing and visualizations in Python contributing to the democratization of scientific tools and techniques and spurring innovation and scientific discovery [100].

Specifically for geoscience, the *Pangeo* project [86] has been a key driver in promoting the use of open-source, high-quality scientific software for big data geoscience that allow researchers to handle petabyte-scale datasets on HPC and cloud platforms. Among others, two key cornerstones of the *Pangeo* software ecosystem are the open sources Python packages *xarray* [53] and *dask* [24]. On the one hand side, *xarray* introduces labels in the form of dimensions, coordinates, and attributes on top of raw NumPy-like arrays, which allows for more intuitive, more concise, and less error-prone user experience. On the other hand side, *dask* facilitates parallel and distributed out-of-memory computing of datasets. Both tools in combination facilitate the analysis of large climate datasets by providing a scalable, efficient and simple way to handle the vast amounts of data produced by climate models, enabling researchers to focus on their scientific questions rather than technical challenges.

2.2 Technical Barriers in Large-Scale and Multi-Dimensional Data Analysis

Climate data analysis often involves the application of multivariate statistical analysis and dimensionality reduction techniques to extract meaningful patterns from large labeled datasets. In Chapter 1 we only presented a tiny fraction of the vast universe of statistical methods for pattern identification in environment and climate science [43, 54]. Commonly, these methods rely on a fundamental assumption: data can be represented as a compilation of observations (or samples) of certain variables (or features). This forms a 2D matrix. Yet, when examining data commonly encountered in Earth observations, be it climate simulations or satellite data, two characteristics of the data stand out:

1. Multi-dimensional Nature: The data spans multiple dimensions, including spatial coordinates (longitude, latitude, height), time metrics (time, steps, forecast lead times), and other categories (variable types, sensors).

2. Large Volume: These datasets are often enormous, frequently exceeding the memory capacity of a single computer.

These two aspects have probably been the key drivers for the de-facto adoption of xarray and dask as the standard in the Python ecosystem when it comes to handling large, multi-dimensional data such as climate and Earth observations.

However, applying dimensionality reduction techniques like the ones discussed in Section 1 can pose challenges to researchers. The typical data processing pipeline for a climate data set (with dimensions time longitude and latitude) includes applying the desired variable weighting, removing invalid grid points (e.g. sea surface temperature with land grid cells), transforming the multidimensional dataset into a 2D matrix format using xarray which is then fed into packages like scikit-learn or dask for out-of-memory computations. These factors can make the analysis of large data sets cumbersome for the researchers. For instance, if one is interested in producing an approximate reconstruction of the original datafield using the first M modes of PCA, after obtaining the 2D reconstruction (e.g. from scikit-learn) one has to reshape the data into its original in its multidimensional form, reinsert land grid cells and unweight the data. While manageable, this process is cumbersome, error-prone, and time-consuming task.

Earlier efforts in this regard have been the development of the python libraries *eoifs* [25] and *xMCA* [47] which addressed PCA and MCA in particular. However, these libraries exhibit limitations when it comes to xarray and dask compatibility, which are essential for handling large climate datasets. For example, there is no support for xarray's fundamental Dataset object and neither of both makes use of randomized linear algebra [33, 42] to greatly speed up the computations.

In addition, climate scientist often require more than just the dimensionality reduction itself. They need to assess the significance of the extracted patterns, e.g. by performing a bootstrap analysis to determine the significance of the modes. This requires additional code and expertise, which can be a barrier to entry for many researchers.

And finally, many classical and traditional methods developed over the last decades in and for climate science to extract meaningful patterns to the best of our knowledge have not been implemented in Python/xarray. This can make it difficult for researchers to apply these methods to their data, limiting the range of techniques available to them.

2.3 *xeofs*: Extracting Meaningful Patterns from Climate Data

Here we briefly introduce the *xeofs* library [96], a Python package that aims to facilitate the application of scalable dimensionality reduction techniques in climate science by providing a user-friendly interface to traditional and state-of-the-art methods specific to the climate science

context. It provides a collection of widely-used dimensionality reduction techniques, focusing on methods prevalent in climate sciences. At the same time, the package development follows the common rules of scientific programming [2], clean coding and test design [56], version control [88], robustness [103], usability [73], documentation [66] and common rules of open software development [90]. Specifically for Python, we *xeofs* largely follows the guidelines and best practices presented by the Pangeo project [86] and *pyOpenSci* Packaging Guide [110].

The three key features of the *xeofs* library are:

Multi-Dimensional & Labeled Data Designed for xarray objects, *xeofs* applies dimensionality reduction to multi-dimensional data while maintaining data labels being based on xarray's DataArray and Dataset objects.

Out-of-Memory Computations Supports large datasets via dask xarray objects, allowing researchers to analyze data that exceeds the memory capacity of a single computer and facilitates scaling up to terabyte computations on HPCs.

Computationally Efficient Ensures computational efficiency, especially with large datasets through randomized linear algebra, which greatly speeds up the computations. A experimental comparison of the performance of the *xeofs* with the *eofs* library for conducting PCA on various data sizes shows that the *xeofs* library performance is about one order of magnitude faster for moderate to larger data sets (Fig. 2.2).

Secondary, additional nice-to-have features of the *xeofs* library include:

Handling Missing Values Can manage common cases of missing values such as masked ocean or land cells.

Weighted Analysis Supports weighted analysis, allowing researchers to apply different weights to the data before performing dimensionality reduction. Binary TrueFalse switches for common weighting schemes (i.e. standard deviation, cosine of latitude; see Section 1.1) are provided.

Significance Testing Provides a user-friendly interface for significance testing using bootstrapping [31].

Model Serialization Allows users to save and load trained models in HDF5, Climate and Forecast (CF) convention-compatible netCDF4 [45] or cloud optimized Zarr formats [79].

All this is written in a modular, object-oriented way to make the package easily extendable to implement and include new dimensionality reduction methods.

All this is offered via an easy-to-use API that allows researchers to apply dimensionality reduction techniques to their climate data with minimal effort. The *xeofs* library is available on the Python Package Index (PyPI) and can be installed via pip. The source code is hosted

Table 2.1 — Even better looking table using booktabs

	Method	Reference
Single-Set Analysis	PCA	[52, 74, 87]
	Complex PCA	0.56
	Hilbert PCA	0.63
	Rotated PCA	0.02
	Extended PCA	0.02
	GWPCA	0.02
	Sparse PCA	0.09
	POP	0.05
	OPA	0.05
Cross-Set Analysis	CCA	0.05
	RDA	0.05
	MCA	0.05
	CPCCA	0.05
	Complex CCA	0.05
	Complex RDA	0.05
	Complex MCA	0.05
	Hilbert CCA	0.05
	Hilbert RDA	0.05
	Hilbert MCA	0.05
	Rotated MCA	0.05
	Rotated CPCCA	0.05
Multi-Set Analysis	CCA	0.05
	Regularized CCA	0.05


1 on GitHub, where users can contribute to the development of the library and report issues or
2 bugs. The library is licensed under the MIT license, making it free to use and modify for both
3 academic and commercial purposes.

4 nicrie: combinations are sometimes also possible, like complex rotated PCA, or hilbert rotated PCA

5 nicrie: some methods are in climate science better known under the name using EOF analysis instead of PCA, such as Extended EOF analysis

6 nicrie: rotated methods include varimax and promax rotation schemes

7 nicrie: mention cca-zoo and ristretto library

Figure 2.2 — Performance of Python library *xeofs*.

¹ **outlook**

- ² implement additional state of the art methods [[13](#), [41](#)] and more bootstrapping methods for
- ³ e.g. MCA and rotated PCA [[105](#)].

References

1. Augspurger, T., Malliaraki, E. & Hopkins, J. *Open Source in Environmental Sustainability* (2023) (cit. on p. 24). 2 3
2. Balaban, G., Grytten, I., Rand, K. D., Scheffer, L. & Sandve, G. K. Ten Simple Rules for Quick and Dirty Scientific Programming. *PLOS Computational Biology* **17**, e1008549. doi:10.1371/journal.pcbi.1008549 (2021) (cit. on p. 27). 4 5 6
3. Barnett, T. P. Interaction of the Monsoon and Pacific Trade Wind System at Interannual Time Scales Part I: The Equatorial Zone. *Monthly Weather Review* **111**, 756–773. doi:10.1175/1520-0493(1983)111<0756:IOTMAP>2.0.CO;2 (1983) (cit. on pp. 6, 13). 7 8 9
4. Barnston, A. G. & Livezey, R. E. Classification, Seasonality and Persistence of Low-Frequency Atmospheric Circulation Patterns. *Monthly weather review* **115**, 1083–1126 (1987) (cit. on p. 10). 10 11 12
5. Barreto, N. J. C. *et al.* Maximum Covariance Analysis to Identify Intraseasonal Oscillations over Tropical Brazil. *Climate Dynamics* **49**, 1583–1596. doi:10.1007/s00382-016-3401-3 (2017) (cit. on p. 17). 13 14 15
6. Bishop, C. M. *Pattern Recognition and Machine Learning* 4 (Springer, 2006) (cit. on pp. 1, 3, 7). 16 17
7. Bloomfield, P. & Davis, J. M. Orthogonal Rotation of Complex Principal Components. *International Journal of Climatology* **14**, 759–775. doi:10.1002/joc.3370140706 (1994) (cit. on p. 14). 18 19 20
8. Boashash, B. Estimating and Interpreting the Instantaneous Frequency of a Signal. I. Fundamentals. *Proceedings of the IEEE* **80**, 520–538. doi:10.1109/5.135376 (1992) (cit. on p. 14). 21 22 23
9. Borga, M., Landelius, T. & Knutsson, H. *A Unified Approach to PCA, PLS, MLR and CCA* (Linköping University, Department of Electrical Engineering, 1997) (cit. on p. 18). 24 25
10. Bretherton, C., Smith, C. & Wallace, J. M. An Intercomparison of Methods for Finding Coupled Patterns in Climate Data. *Journal of climate* **5**, 541–560 (1992) (cit. on pp. 16, 17). 26 27
11. Broomhead, D. S. & King, G. P. Extracting Qualitative Dynamics from Experimental Data. *Physica D: Nonlinear Phenomena* **20**, 217–236. doi:10.1016/0167-2789(86)90031-X (1986) (cit. on p. 6). 28 29 30
12. Buell, C. E. *On the Physical Interpretation of Empirical Orthogonal Functions* in. Preprints Sixth Conference on Probability and Statistics in Atmospheric Science (American Meteorological Society, Banff, Alberta, 1979) (cit. on p. 8). 31 32 33
13. Bueso, D., Piles, M. & Camps-Valls, G. Nonlinear PCA for Spatio-Temporal Analysis of Earth Observation Data. *IEEE Transactions on Geoscience and Remote Sensing*, 1–12. doi:10.1109/TGRS.2020.2969813 (2020) (cit. on pp. 6, 11, 30). 34 35 36
14. Bürger, G. Complex Principal Oscillation Pattern Analysis. *Journal of climate* **6**, 1972–1986 (1993) (cit. on p. 14). 37 38
15. Bürger, G. Intraseasonal Oscillation Indices from Complex EOFs. *Journal of Climate* **34**, 107–122. doi:10.1175/JCLI-D-20-0427.1 (2021) (cit. on pp. 13, 14). 39 40

16. C3S. *ERA5 Hourly Data on Single Levels from 1940 to Present* (Copernicus Climate Change Service (C3S) Climate Data Store (CDS), 2018). doi:[10.24381/CDS.ADBB2D47](https://doi.org/10.24381/CDS.ADBB2D47) (cit. on p. 5).
17. Cai, W. & Baines, P. G. Forcing of the Antarctic Circumpolar Wave by El Niño-Southern Oscillation Teleconnections. *Journal of Geophysical Research: Oceans* **106**, 9019–9038. doi:[10.1029/2000JC000590](https://doi.org/10.1029/2000JC000590) (2001) (cit. on p. 13).
18. Camacho, J., Smilde, A. K., Saccenti, E. & Westerhuis, J. A. All Sparse PCA Models Are Wrong, but Some Are Useful. Part I: Computation of Scores, Residuals and Explained Variance. *Chemometrics and Intelligent Laboratory Systems* **196**, 103907 (2020) (cit. on p. 11).
19. Candès, E. J., Li, X., Ma, Y. & Wright, J. Robust Principal Component Analysis? *Journal of the ACM* **58**, 11:1–11:37. doi:[10.1145/1970392.1970395](https://doi.org/10.1145/1970392.1970395) (2011) (cit. on p. 6).
20. Cheng, X. & Dunkerton, T. J. Orthogonal Rotation of Spatial Patterns Derived from Singular Value Decomposition Analysis. *Journal of Climate* **8**, 2631–2643. doi:[10.1175/1520-0442\(1995\)008<2631:OROSPD>2.0.CO;2](https://doi.org/10.1175/1520-0442(1995)008<2631:OROSPD>2.0.CO;2) (1995) (cit. on p. 17).
21. Cheng, X., Nitsche, G. & Wallace, J. M. Robustness of Low-Frequency Circulation Patterns Derived from EOF and Rotated EOF Analyses. *Journal of Climate* **8**, 1709–1713. doi:[10.1175/1520-0442\(1995\)008<1709:ROLFCP>2.0.CO;2](https://doi.org/10.1175/1520-0442(1995)008<1709:ROLFCP>2.0.CO;2) (1995) (cit. on p. 10).
22. Cherry, S. Singular Value Decomposition Analysis and Canonical Correlation Analysis. *Journal of Climate* **9**, 2003–2009 (1996) (cit. on p. 17).
23. Cherry, S. Some Comments on Singular Value Decomposition Analysis. *Journal of Climate* **10**, 1759–1761 (1997) (cit. on p. 18).
24. Dask Development Team. *Dask: Library for Dynamic Task Scheduling* manual (2016) (cit. on p. 25).
25. Dawson, A. Eofs: A Library for EOF Analysis of Meteorological, Oceanographic, and Climate Data. **4**, e14. doi:[10.5334/jors.122](https://doi.org/10.5334/jors.122) (1 2016) (cit. on p. 26).
26. DelSole, T. Optimally Persistent Patterns in Time-Varying Fields. *Journal of the Atmospheric Sciences* **58**, 1341–1356. doi:[10.1175/1520-0469\(2001\)058<1341:OPPITV>2.0.CO;2](https://doi.org/10.1175/1520-0469(2001)058<1341:OPPITV>2.0.CO;2) (2001) (cit. on p. 6).
27. DelSole, T. Low-Frequency Variations of Surface Temperature in Observations and Simulations. *Journal of Climate* **19**, 4487–4507. doi:[10.1175/JCLI3879.1](https://doi.org/10.1175/JCLI3879.1) (2006) (cit. on p. 6).
28. Delsole, T. M. & Tippett, M. K. *Statistical Methods for Climate Scientists* (Cambridge University Press, New York, NY, USA, 2022) (cit. on p. 1).
29. Drineas, P., Frieze, A., Kannan, R., Vempala, S. & Vinay, V. Clustering Large Graphs via the Singular Value Decomposition. *Machine Learning* **56**, 9–33. doi:[10.1023/B:MACH.0000033113.59016.96](https://doi.org/10.1023/B:MACH.0000033113.59016.96) (2004) (cit. on p. 5).
30. Easterbrook, S. M. & Johns, T. C. Engineering the Software for Understanding Climate Change. *Computing in science & engineering* **11**, 65–74 (2009) (cit. on p. 24).
31. Efron, B. Bootstrap Methods: Another Look at the Jackknife. *The Annals of Statistics* **7**. doi:[10.1214/aos/1176344552](https://doi.org/10.1214/aos/1176344552) (1979) (cit. on p. 27).
32. Elipot, S., Frajka-Williams, E., Hughes, C. W., Olhede, S. & Lankhorst, M. Observed Basin-Scale Response of the North Atlantic Meridional Overturning Circulation to Wind Stress Forcing. *Journal of Climate* **30**, 2029–2054. doi:[10.1175/JCLI-D-16-0664.1](https://doi.org/10.1175/JCLI-D-16-0664.1) (2017) (cit. on p. 18).
33. Erichson, N. B., Voronin, S., Brunton, S. L. & Kutz, J. N. Randomized Matrix Decompositions Using R. *Journal of Statistical Software* **89**. doi:[10.18637/jss.v089.i11](https://doi.org/10.18637/jss.v089.i11) (2019) (cit. on p. 26).

34. Erichson, N. B. *et al.* Sparse Principal Component Analysis via Variable Projection. *SIAM Journal on Applied Mathematics* **80**, 977–1002. doi:[10.1137/18M1211350](#) (2020) (cit. on p. [11](#)). 1
2
3
35. Farrell, B. F. & Ioannou, P. J. Accurate Low-Dimensional Approximation of the Linear Dynamics of Fluid Flow (2001) (cit. on p. [6](#)). 4
5
36. Frankignoul, C. & Kestenare, E. Air–Sea Interactions in the Tropical Atlantic: A View Based on Lagged Rotated Maximum Covariance Analysis. *Journal of Climate* **18**, 3874–3890. doi:[10.1175/JCLI3498.1](#) (2005) (cit. on p. [17](#)). 6
7
8
37. Gabor, D. Theory of Communication. Part 1: The Analysis of Information. *Journal of the Institution of Electrical Engineers - Part III: Radio and Communication Engineering* **93**, 429–441. doi:[10.1049/ji-3-2.1946.0074](#) (1946) (cit. on p. [14](#)). 9
10
11
38. Gao, T., Yu, J.-y. & Paek, H. Impacts of Four Northern-Hemisphere Teleconnection Patterns on Atmospheric Circulations over Eurasia and the Pacific. *Theoretical and Applied Climatology* **129**, 815–831 (2017) (cit. on p. [10](#)). 12
13
14
39. Golub, G. H. & Van Loan, C. F. *Matrix Computations* (JHU press, 2013) (cit. on p. [5](#)). 15
40. Guerra-Urzola, R., Van Deun, K., Vera, J. C. & Sijtsma, K. A Guide for Sparse PCA: Model Comparison and Applications. *Psychometrika* **86**, 893–919. doi:[10.1007/s11336-021-09773-2](#) (2021) (cit. on p. [11](#)). 16
17
18
41. Guilloteau, C., Mamalakis, A., Vulis, L., Georgiou, T. T. & Foufoula-Georgiou, E. *Rotated Spectral Principal Component Analysis (rsPCA) for Identifying Dynamical Modes of Variability in Climate Systems* 2020 (cit. on pp. [11](#), [30](#)). 19
20
21
42. Halko, N., Martinsson, P. G. & Tropp, J. A. Finding Structure with Randomness: Probabilistic Algorithms for Constructing Approximate Matrix Decompositions. *SIAM Review* **53**, 217–288. doi:[10.1137/090771806](#) (2011) (cit. on pp. [5](#), [7](#), [26](#)). 22
23
24
43. Hannachi, A. *Patterns Identification and Data Mining in Weather and Climate* 600 pp. (Springer International Publishing, 2021) (cit. on pp. [1](#), [13](#), [17](#), [25](#)). 25
26
44. Harris, C. R. *et al.* Array Programming with NumPy. *Nature* **585**, 357–362. doi:[10.1038/s41586-020-2649-2](#) (2020) (cit. on p. [25](#)). 27
28
45. Hassell, D., Gregory, J., Blower, J., Lawrence, B. N. & Taylor, K. E. A Data Model of the Climate and Forecast Metadata Conventions (CF-1.6) with a Software Implementation (Cf-Python v2.1). *Geoscientific Model Development* **10**, 4619–4646. doi:[10.5194/gmd-10-4619-2017](#) (2017) (cit. on p. [27](#)). 29
30
31
32
46. Hasselmann, K. PIPs and POPs: The Reduction of Complex Dynamical Systems Using Principal Interaction and Oscillation Patterns. *Journal of Geophysical Research: Atmospheres* **93**, 11015–11021. doi:[10.1029/JD093iD09p11015](#) (1988) (cit. on p. [6](#)). 33
34
35
47. He, C. *xMCA: Maximum Covariance Analysis in Xarray for Climate Science* 2020 (cit. on p. [26](#)). 36
37
48. Hendrickson, A. E. & White, P. O. Promax: A Quick Method for Rotation to Oblique Simple Structure. *British Journal of Statistical Psychology* **17**, 65–70. doi:[10.1111/j.2044-8317.1964.tb00244.x](#) (1964) (cit. on p. [9](#)). 38
39
40
49. Horel, J. D. Complex Principal Component Analysis: Theory and Examples. *Journal of Climate and Applied Meteorology* **23**, 1660–1673. doi:[10.1175/1520-0450\(1984\)023<1660:CPCATA>2.0.CO;2](#) (1984) (cit. on pp. [6](#), [11](#), [13](#)). 41
42
43
50. Horel, J. D. A Rotated Principal Component Analysis of the Interannual Variability of the Northern Hemisphere 500 Mb Height Field. *Monthly Weather Review* **109**, 2080–2092 (1981) (cit. on pp. [8](#), [10](#), [14](#)). 44
45
46

- 1 51. Horn, R. A. & Johnson, C. R. *Matrix Analysis* (Cambridge university press, 2012) (cit. on
2 pp. 3, 4).
- 3 52. Hotelling, H. Relations Between Two Sets of Variates. *Biometrika* **28**, 321–377. doi:10.
4 2307/2333955 (1936) (cit. on pp. 1, 28).
- 5 53. Hoyer, S. & Hamman, J. Xarray: N-D Labeled Arrays and Datasets in Python. *Journal of*
6 *Open Research Software* **5**, 10. doi:10.5334/jors.148 (1 2017) (cit. on p. 25).
- 7 54. Hsieh, W. W. *Introduction to Environmental Data Science* (Cambridge University Press,
8 New York, 2023) (cit. on p. 25).
- 9 55. Hunter, J. D. Matplotlib: A 2D Graphics Environment. *Computing in Science & Engineering*
10 **9**, 90–95. doi:10.1109/MCSE.2007.55 (2007) (cit. on p. 25).
- 11 56. Hunter-Zinck, H., Siqueira, A. F. de, Vasquez, V. N., Barnes, R. & Martinez, C. C. Ten
12 Simple Rules on Writing Clean and Reliable Open-Source Scientific Software. *PLoS*
13 *Computational Biology* **17**. doi:10.1371/journal.pcbi.1009481 (2021) (cit. on p. 27).
- 14 57. Hurley, J. R. & Cattell, R. B. The Procrustes Program: Producing Direct Rotation to Test a
15 Hypothesized Factor Structure. *Behavioral science* **7**, 258 (1962) (cit. on p. 9).
- 16 58. Huth, R. & Beranová, R. How to Recognize a True Mode of Atmospheric Circulation
17 Variability. *Earth and Space Science* **8**, e2020EA001275. doi:10.1029/2020EA001275 (2021)
18 (cit. on p. 10).
- 19 59. Ince, D. C., Hatton, L. & Graham-Cumming, J. The Case for Open Computer Programs.
20 *Nature* **482**, 485–488. doi:10.1038/nature10836 (2012) (cit. on p. 25).
- 21 60. Jolliffe, I. T. *Principal Component Analysis* 2nd ed. (Springer, New York, 2002) (cit. on
22 pp. 1, 3, 7, 8, 11, 13, 17, 18).
- 23 61. Jolliffe, I., Uddin, M. & Vines, S. Simplified EOFs-three Alternatives to Rotation. *Climate*
24 *Research* **20**, 271–279. doi:10.3354/cr020271 (2002) (cit. on pp. 6, 11).
- 25 62. Kaiser, H. F. The Varimax Criterion for Analytic Rotation in Factor Analysis. *Psychome-*
26 *trika* **23**, 187–200. doi:10.1007/BF02289233 (1958) (cit. on p. 9).
- 27 63. Kettenring, J. R. Canonical Analysis of Several Sets of Variables. *Biometrika* **58**, 433–451.
28 doi:10.2307/2334380 (1971) (cit. on p. 14).
- 29 64. Kramer, M. A. Nonlinear Principal Component Analysis Using Autoassociative Neural
30 Networks. *AIChE Journal* **37**, 233–243. doi:10.1002/aic.690370209 (1991) (cit. on p. 6).
- 31 65. Krogh, G., Spaeth, S. & Lakhani, K. Community, Joining, and Specialization in Open
32 Source Software Innovation: A Case Study. *Research Policy* **32**, 1217–1241. doi:10.1016/
33 S0048-7333(03)00050-7 (2003) (cit. on p. 25).
- 34 66. Lee, B. D. Ten Simple Rules for Documenting Scientific Software. *PLOS Computational*
35 *Biology* **14**, e1006561. doi:10.1371/journal.pcbi.1006561 (2018) (cit. on p. 27).
- 36 67. Lee, J., Sperber, K. R., Gleckler, P. J., Bonfils, C. J. & Taylor, K. E. Quantifying the Agree-
37 ment between Observed and Simulated Extratropical Modes of Interannual Variability.
38 *Climate Dynamics* **52**, 4057–4089 (2019) (cit. on p. 10).
- 39 68. Legates, D. R. The Effect of Domain Shape on Principal Components Analyses. *Inter-*
40 *national Journal of Climatology* **11**, 135–146. doi:10.1002/joc.3370110203 (1991) (cit. on
41 p. 8).
- 42 69. Legates, D. R. The Effect of Domain Shape on Principal Components Analyses: A Reply.
43 *International Journal of Climatology* **13**, 219–228. doi:10.1002/joc.3370130207 (1993) (cit. on
44 p. 7).

70. Li, L., Schmitt, R. W., Ummenhofer, C. C. & Karnauskas, K. B. North Atlantic Salinity as a Predictor of Sahel Rainfall. *Science Advances* **2**, e1501588. doi:[10.1126/sciadv.1501588](https://doi.org/10.1126/sciadv.1501588) (2016) (cit. on p. 17). 1
2
3
71. Lian, T. & Chen, D. An Evaluation of Rotated EOF Analysis and Its Application to Tropical Pacific SST Variability. *Journal of Climate* **25**, 5361–5373. doi:[10.1175/JCLI-D-11-00663.1](https://doi.org/10.1175/JCLI-D-11-00663.1) (2012) (cit. on p. 10). 4
5
6
72. Lim, Y.-K. The East Atlantic/West Russia (EA/WR) Teleconnection in the North Atlantic: Climate Impact and Relation to Rossby Wave Propagation. *Climate dynamics* **44**, 3211–3222 (2015) (cit. on p. 10). 7
8
9
73. List, M., Ebert, P. & Albrecht, F. Ten Simple Rules for Developing Usable Software in Computational Biology. *PLOS Computational Biology* **13**, e1005265. doi:[10.1371/journal.pcbi.1005265](https://doi.org/10.1371/journal.pcbi.1005265) (2017) (cit. on p. 27). 10
11
12
74. Lorenz, E. N. *Empirical Orthogonal Functions and Statistical Weather Prediction* Scientific Report 1 (Massachusetts Institute of Technology, Department of Meteorology, Cambridge, 1956) (cit. on pp. 1, 28). 13
14
15
75. Madden, R. A. & Julian, P. R. Detection of a 40–50 Day Oscillation in the Zonal Wind in the Tropical Pacific. *Journal of the Atmospheric Sciences* **28**, 702–708. doi:[10.1175/1520-0469\(1971\)028<0702:DOADOI>2.0.CO;2](https://doi.org/10.1175/1520-0469(1971)028<0702:DOADOI>2.0.CO;2) (1971) (cit. on p. 14). 16
17
18
76. Merrifield, M. A. & Guza, R. T. Detecting Propagating Signals with Complex Empirical Orthogonal Functions: A Cautionary Note. *Journal of Physical Oceanography* **20**, 1628–1633. doi:[10.1175/1520-0485\(1990\)020<1628:DPSWCE>2.0.CO;2](https://doi.org/10.1175/1520-0485(1990)020<1628:DPSWCE>2.0.CO;2) (1990) (cit. on p. 13). 19
20
21
77. Mestas-Nuñez, A. M. Orthogonality Properties of Rotated Empirical Modes. *International Journal of Climatology* **20**, 1509–1516. doi:[10.1002/1097-0088\(200010\)20:12<1509::AID-JOC553>3.0.CO;2-Q](https://doi.org/10.1002/1097-0088(200010)20:12<1509::AID-JOC553>3.0.CO;2-Q) (2000) (cit. on p. 8). 22
23
24
78. Mezzina, B., García-Serrano, J., Bladé, I. & Kucharski, F. Dynamics of the ENSO Teleconnection and NAO Variability in the North Atlantic–European Late Winter. *Journal of Climate* **33**, 907–923 (2020) (cit. on p. 10). 25
26
27
79. Miles, A. et al. *Zarr-Developers/Zarr-Python: V2.4.0* version v2.4.0. Zenodo, 2020. doi:[10.5281/zenodo.3773450](https://doi.org/10.5281/zenodo.3773450) (cit. on p. 27). 28
29
80. Millman, K. J. & Pérez, F. in *Implementing Reproducible Research* 149–183 (Chapman and Hall/CRC, 2018) (cit. on p. 25). 30
31
81. Mori, M., Kosaka, Y., Watanabe, M., Nakamura, H. & Kimoto, M. A Reconciled Estimate of the Influence of Arctic Sea-Ice Loss on Recent Eurasian Cooling. *Nature Climate Change* **9**, 123–129. doi:[10.1038/s41558-018-0379-3](https://doi.org/10.1038/s41558-018-0379-3) (2019) (cit. on p. 17). 32
33
34
82. North, G. R., Moeng, F. J., Bell, T. L. & Cahalan, R. F. The Latitude Dependence of the Variance of Zonally Averaged Quantities. *Monthly Weather Review* **110**, 319–326 (1982) (cit. on p. 7). 35
36
37
83. Oliphant, T. E. Python for Scientific Computing. *Computing in science & engineering* **9**, 10–20 (2007) (cit. on p. 25). 38
39
84. Oppenheim, A. V. & Schafer, R. W. *Discrete-Time Signal Processing* 3rd (Pearson Higher Education, 2010) (cit. on p. 13). 40
41
85. Overpeck, J. T., Meehl, G. A., Bony, S. & Easterling, D. R. Climate Data Challenges in the 21st Century. *Science* **331**, 700–702. doi:[10.1126/science.1197869](https://doi.org/10.1126/science.1197869) (2011) (cit. on p. 23). 42
43
86. *Pangeo - A Community Platform for Big Data Geoscience* <https://pangeo.io/index.html> (2024) (cit. on pp. 25, 27). 44
45

- 1 87. Pearson, K. On Lines and Planes of Closest Fit to Systems of Points in Space. *The*
2 *London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science* **2**, 559–572.
3 doi:[10.1080/14786440109462720](https://doi.org/10.1080/14786440109462720) (1901) (cit. on pp. [1](#), [28](#)).
- 4 88. Perez-Riverol, Y. *et al.* Ten Simple Rules for Taking Advantage of Git and GitHub. *PLOS*
5 *Computational Biology* **12**, e1004947. doi:[10.1371/journal.pcbi.1004947](https://doi.org/10.1371/journal.pcbi.1004947) (2016) (cit. on
6 p. [27](#)).
- 7 89. Preisendorfer, R. W. & Mobley, C. D. *Principal Component Analysis in Meteorology and*
8 *Oceanography* (Elsevier Science Publishers, 1988) (cit. on p. [1](#)).
- 9 90. Prlić, A. & Procter, J. B. Ten Simple Rules for the Open Development of Scientific Software.
10 *PLOS Computational Biology* **8**, e1002802. doi:[10.1371/journal.pcbi.1002802](https://doi.org/10.1371/journal.pcbi.1002802) (2012) (cit. on
11 p. [27](#)).
- 12 91. Rasmusson, E. M., Arkin, P. A., Chen, W.-Y. & Jalickee, J. B. Biennial Variations in Surface
13 Temperature over the United States as Revealed by Singular Decomposition. *Monthly*
14 *Weather Review* **109**, 587–598 (1981) (cit. on p. [13](#)).
- 15 92. Rathore, S. *et al.* Improving Australian Rainfall Prediction Using Sea Surface Salinity.
16 *Journal of Climate* **34**, 2473–2490 (2021) (cit. on p. [17](#)).
- 17 93. Richman, M. B. Obliquely Rotated Principal Components: An Improved Meteorological
18 Map Typing Technique? *Journal of Applied Meteorology and Climatology* **20**, 1145–1159
19 (1981) (cit. on p. [8](#)).
- 20 94. Richman, M. B. Rotation of Principal Components. *Journal of Climatology* **6**, 293–335.
21 doi:[10.1002/joc.3370060305](https://doi.org/10.1002/joc.3370060305) (1986) (cit. on pp. [6–10](#)).
- 22 95. Richman, M. B. & Lamb, P. J. Climatic Pattern Analysis of Three-and Seven-Day Summer
23 Rainfall in the Central United States: Some Methodological Considerations and a Re-
24 gionalization. *Journal of Applied Meteorology and Climatology* **24**, 1325–1343 (1985) (cit. on
25 p. [10](#)).
- 26 96. Rieger, N. & Levang, S. J. xeoofs: Comprehensive EOF Analysis in Python with Xarray.
27 *Journal of Open Source Software* **9**, 6060. doi:[10.21105/joss.06060](https://doi.org/10.21105/joss.06060) (2024) (cit. on pp. [23](#), [26](#)).
- 28 97. Sampson, P. D., Streissguth, A. P., Barr, H. M. & Bookstein, F. L. Neurobehavioral Effects
29 of Prenatal Alcohol: Part II. Partial Least Squares Analysis. *Neurotoxicology and teratology*
30 **11**, 477–491 (1989) (cit. on p. [17](#)).
- 31 98. Schnase, J. L. *et al.* Big Data Challenges in Climate Science. *IEEE geoscience and remote*
32 *sensing magazine* **Volume 4**, 10–22. doi:[10.1109/MGRS.2015.2514192](https://doi.org/10.1109/MGRS.2015.2514192) (Iss 3 2016) (cit. on
33 p. [24](#)).
- 34 99. Schölkopf, B., Smola, A. & Müller, K.-R. Nonlinear Component Analysis as a Kernel
35 Eigenvalue Problem. *Neural Computation* **10**, 1299–1319. doi:[10.1162/089976698300017467](https://doi.org/10.1162/089976698300017467)
36 (1998) (cit. on p. [6](#)).
- 37 100. Sonnenburg, S. *et al.* The Need for Open Source Software in Machine Learning. *Journal*
38 *of Machine Learning Research* **8**, 2443–2466. doi:[10.5555/1314498.1314577](https://doi.org/10.5555/1314498.1314577) (2007) (cit. on
39 p. [25](#)).
- 40 101. Storch, H. von & Zwiers, F. W. *Statistical Analysis in Climate Research* (Cambridge Univer-
41 sity Press (Virtual Publishing), 2003) (cit. on pp. [1](#), [3](#), [12–14](#)).
- 42 102. Swenson, E. Continuum Power CCA: A Unified Approach for Isolating Coupled Modes.
43 *Journal of Climate* **28**, 1016–1030. doi:[10.1175/JCLI-D-14-00451.1](https://doi.org/10.1175/JCLI-D-14-00451.1) (2015) (cit. on p. [18](#)).
- 44 103. Taschuk, M. & Wilson, G. Ten Simple Rules for Making Research Software More Robust.
45 *PLOS Computational Biology* **13**, e1005412. doi:[10.1371/journal.pcbi.1005412](https://doi.org/10.1371/journal.pcbi.1005412) (2017) (cit. on
46 p. [27](#)).

104. Tibshirani, R. Regression Shrinkage and Selection via the Lasso. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **58**, 267–288 (1996) (cit. on p. 11). 1
2
105. Timmerman, M. E., Kiers, H. A. L. & Smilde, A. K. Estimating Confidence Intervals for Principal Component Loadings: A Comparison between the Bootstrap and Asymptotic Results. *British Journal of Mathematical and Statistical Psychology* **60**, 295–314. doi:[10.1348/000711006X109636](https://doi.org/10.1348/000711006X109636) (2007) (cit. on p. 30). 3
4
5
6
106. Trendafilov, N. T. & Jolliffe, I. T. Projected Gradient Approach to the Numerical Solution of the SCoTLASS. *Computational Statistics & Data Analysis* **50**, 242–253 (2006) (cit. on p. 11). 7
8
9
107. Tsai, H.-C. & Lee, T. H. Maximum Covariance Analysis of Typhoon Surface Wind and Rainfall Relationships in Taiwan. *Journal of applied meteorology and climatology* **48**, 997–1016 (2009) (cit. on p. 17). 10
11
12
108. Virtanen, P. *et al.* SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods* **17**, 261–272. doi:[10.1038/s41592-019-0686-2](https://doi.org/10.1038/s41592-019-0686-2) (2020) (cit. on p. 25). 13
14
109. Wallace, J. M., Smith, C. & Bretherton, C. S. Singular Value Decomposition of Wintertime Sea Surface Temperature and 500-Mb Height Anomalies. *Journal of climate* **5**, 561–576 (1992) (cit. on p. 17). 15
16
17
110. Wasser, L. *et al.* *pyOpenSci/Python-Package-Guide: pyOpenSci Packaging Guide v0.4* doi:[10.5281/zenodo.13154979](https://doi.org/10.5281/zenodo.13154979) (Zenodo, 2024) (cit. on p. 27). 18
19
111. Weare, B. C. & Nasstrom, J. S. Examples of Extended Empirical Orthogonal Function Analyses. *Monthly Weather Review* **110**, 481–485. doi:[10.1175/1520-0493\(1982\)110<0481:EOEEOF>2.0.CO;2](https://doi.org/10.1175/1520-0493(1982)110<0481:EOEEOF>2.0.CO;2) (1982) (cit. on p. 6). 20
21
22
112. Webb, A. R. An Approach to Non-Linear Principal Components Analysis Using Radially Symmetric Kernel Functions. *Statistics and Computing* **6**, 159–168. doi:[10.1007/BF00162527](https://doi.org/10.1007/BF00162527) (1996) (cit. on p. 6). 23
24
25
113. Wegelin, J. A. *A Survey of Partial Least Squares (PLS) Methods, with Emphasis on the Two-Block Case* 371 (University of Washington, Seattle, Washington, 98195, USA, 2000) (cit. on p. 17). 26
27
28
114. Wilks, D. S. *Statistical Methods in the Atmospheric Sciences* Fourth Edition. 698 pp. doi:[10.1016/B978-0-12-815823-4.00011-0](https://doi.org/10.1016/B978-0-12-815823-4.00011-0) (Academic Press, 2019) (cit. on pp. 1, 18). 29
30
115. Zou, H., Hastie, T. & Tibshirani, R. Sparse Principal Component Analysis. *Journal of Computational and Graphical Statistics*. doi:[10.1198/106186006X113430](https://doi.org/10.1198/106186006X113430) (2006) (cit. on p. 11). 31
32
33
116. Zou, H. & Xue, L. A Selective Overview of Sparse Principal Component Analysis. *Proceedings of the IEEE* **106**, 1311–1320 (2018) (cit. on p. 11). 34
35



Proofs

The Frobenius norm $\|\cdot\|_F$ of a complex $N \times P$ matrix $\mathbf{A} \in \mathbb{C}^{N \times P}$ is defined as:

$$\|\mathbf{A}\|_F \equiv \left(\sum_i^N \sum_j^P |\alpha_{ij}|^2 \right)^{1/2} = \left(\sum_i^{\min(N,P)} \sigma_i^2(\mathbf{A}) \right)^{1/2}. \quad (\text{A.1})$$

This expression reveals that the Frobenius norm is equivalent to the ℓ_2 -norm of the matrix's singular values σ_i . We can demonstrate this equivalence using the singular value decomposition (SVD) of $\mathbf{A}$ and leveraging the property that the trace of a matrix remains invariant under cyclic permutations.

Starting with the definition of the Frobenius norm, we have:

$$\|\mathbf{A}\|_F = \left(\text{tr}(\mathbf{A}^H \mathbf{A}) \right)^{1/2}$$

Applying the SVD, we can express $\mathbf{A} = \mathbf{V} \mathbf{\Sigma} \mathbf{V}^H$, leading to:

$$\|\mathbf{A}\|_F = \left(\text{tr}(\mathbf{V} \mathbf{\Sigma} \mathbf{U}^H \mathbf{U} \mathbf{\Sigma}^H \mathbf{V}^H) \right)^{1/2}$$

Since $\mathbf{U}^H \mathbf{U} = \mathbf{I}$, this simplifies to:

$$\|\mathbf{A}\|_F = \left(\text{tr}(\mathbf{V} \mathbf{\Sigma}^2 \mathbf{V}^H) \right)^{1/2}$$

Finally using the invariance property of the trace, we obtain:

1

$$\|\mathbf{A}\|_{\text{F}} = \left(\sum_i^{\min(N,P)} \sigma_i^2(\mathbf{A}) \right)^{1/2}$$

- 2 This concludes the proof that the Frobenius norm of a matrix is indeed the ℓ_2 -norm of its
3 singular values.

